

# Hybrid RAG System - Final Report

**Group Number:** 94 **Date:** 7th February 2026 **GitHub Repository:** DarshanPatel11/CI\_Assignment\_2

**Note:** For the most accurate evaluation metrics, please refer to the latest evaluation reports in `data/evaluation/results/`. Earlier reports may show lower performance due to ongoing testing and bug fixes during development.

---

## Table of Contents

1. Executive Summary
  2. System Architecture
  3. Data Pipeline
  4. Evaluation Framework
  5. Evaluation Metrics
  6. Results
  7. Ablation Study
  8. Error Analysis
  9. User Interface
  10. Innovative Approaches
  11. Deployment
  12. Conclusion
  13. References
- 

## 1. Executive Summary

This report documents the implementation of a **Hybrid Retrieval-Augmented Generation (RAG) system** that combines:

- Dense Vector Retrieval (FAISS with all-MiniLM-L6-v2)
- Sparse Keyword Retrieval (BM25)
- Reciprocal Rank Fusion (RRF) for optimal result merging
- Flan-T5-base for answer generation

## Key Results Summary

Metric	Score	Interpretation
<b>MRR (URL-level)</b>	<b>0.913</b>	Excellent source retrieval
<b>Hit Rate</b>	<b>95%</b>	Very high recall
<b>Faithfulness</b>	0.32	Room for improvement
<b>Context Precision</b>	0.68	Effective ranking

Metric	Score	Interpretation
<b>Mean Response Time</b>	~250ms	Fast response

## 2. System Architecture

### 2.1 Architecture Diagrams




Figure 1: Hybrid RAG System Architecture

### Hybrid RAG System

#### Data Pipeline

#### Evaluation Pipeline

### 2.2 Component Summary

Component	Technology	Purpose
<b>Embeddings</b>	all-MiniLM-L6-v2	384-dim semantic vectors
<b>Vector Index</b>	FAISS (IndexFlatIP)	Cosine similarity search

Component	Technology	Purpose
<b>Sparse Search</b>	BM25Okapi	Keyword-based retrieval
<b>Fusion</b>	RRF (k=60)	Combines rankings
<b>LLM</b>	Flan-T5-base	Answer generation
<b>UI</b>	Streamlit	Interactive interface
<b>Deployment</b>	Docker	Containerized setup

### 2.3 Data Flow

1. **Query Input** → Streamlit UI
  2. **Dense Encoding** → all-MiniLM-L6-v2 → FAISS search
  3. **Sparse Search** → BM25 tokenization → score calculation
  4. **RRF Fusion** → Combine rankings with k=60
  5. **Context Selection** → Top-N chunks (truncated to 450 tokens)
  6. **Generation** → Flan-T5 produces answer
  7. **Response** → Answer + sources + scores
- 

## 3. Data Pipeline

### 3.1 Wikipedia Corpus

Category	Count	Description
<b>Fixed URLs</b>	200	Unique set covering diverse topics
<b>Random URLs</b>	300	Dynamically sampled per indexing run
<b>Total Corpus</b>	500	Combined dataset

### 3.2 Text Processing

- **Minimum article length:** 200 words
- **Chunk size:** 200-400 tokens
- **Overlap:** 50 tokens
- **Metadata preserved:** URL, title, unique chunk ID

### 3.3 Fixed URL Categories

The 200 fixed URLs cover diverse topics including: - Science, Technology, History, Geography, Arts - Philosophy, Literature, Mathematics, Biology, Physics - Chemistry, Medicine, Economics, Politics, Sports - Music, Film, Architecture, Psychology, Sociology

See `data/fixed_urls.json` for the complete list.

---




Figure 2: Data Pipeline




Figure 3: Evaluation Pipeline

## 4. Evaluation Framework

### 4.1 Question Dataset

**Total Questions:** 100 Q&A pairs generated from the Wikipedia corpus

Question Type	Count	Percentage
Factual	51	51%
Comparative	20	20%
Inferential	20	20%
Multi-hop	9	9%

Difficulty	Count	Percentage
Easy	21	21%
Medium	62	62%
Hard	17	17%

### 4.2 Automated Pipeline

The evaluation pipeline runs with a single command:

```
python main.py --evaluate --innovative --num-questions 100
```

This automatically: 1. Loads the 100-question dataset 2. Runs the RAG system on each question 3. Computes all metrics (MRR, Faithfulness, Context Precision) 4. Generates comprehensive reports (JSON, CSV, HTML)

---

## 5. Evaluation Metrics

### 5.1 Mandatory Metric: MRR (Mean Reciprocal Rank)

**URL-Level Evaluation** - Measures how quickly the system finds the correct source document.

#### Calculation Method

$$MRR = \frac{1}{|Q|} \sum_{i=1}^{|Q|} \frac{1}{rank_i}$$

Where: -  $|Q|$  = Total number of queries -  $rank_i$  = Position of the first correct Wikipedia URL in retrieved results

## Interpretation

MRR Score	Quality
0.9 - 1.0	Excellent
0.7 - 0.9	Good
0.5 - 0.7	Fair
< 0.5	Poor

**Our Result:** MRR = **0.913** (Excellent)

---

## 5.2 Custom Metric 1: Faithfulness Score (LLM-as-Judge)

**Justification** Faithfulness is critical for RAG systems because:

- **Hallucination Detection:** Identifies when the model generates information not present in retrieved context
- **Trust:** Users need to trust that answers are grounded in actual sources
- **Quality Assurance:** Ensures the system doesn't "make things up"

### Calculation Method

1. **Claim Extraction:** Use LLM to extract atomic claims from the generated answer
2. **Claim Verification:** For each claim, verify if it's supported by the retrieved context
3. **Score Calculation:**

$$Faithfulness = \frac{\text{Number of Supported Claims}}{\text{Total Number of Claims}}$$

For each claim, verification uses:  
- LLM-based verification prompting  
- Fallback semantic similarity (threshold > 0.7)

### Implementation Details

```
def calculate_faithfulness_single(self, answer: str, context: str):  
    # Step 1: Extract claims using LLM  
    claims = self._extract_claims(answer)  
  
    # Step 2: Verify each claim against context  
    supported = 0  
    for claim in claims:  
        if self._verify_claim(claim, context):
```

```

    supported += 1

# Step 3: Calculate score
return supported / max(len(claims), 1)

```

### Interpretation

Score	Interpretation
1.0	Fully grounded - all claims supported by context
0.7-0.99	Mostly grounded - some minor unsupported claims
0.5-0.69	Partially grounded - mix of supported/unsupported
< 0.5	Reliability concerns - significant hallucination risk

**Our Result: Faithfulness = 0.32** (Need larger LLM)

---

### 5.3 Custom Metric 2: Context Precision

**Justification** Context Precision is important because: - **Ranking Quality**: Measures if relevant documents appear at top ranks - **Efficiency**: Higher-ranked relevant docs → better context for generation - **Beyond Simple Recall**: Evaluates the ORDER of results, not just presence

**Calculation Method** Context Precision uses weighted precision at each rank position:

$$CP = \frac{\sum_{k=1}^K (Precision@k \times rel_k)}{\text{Total Relevant Documents}}$$

Where: -  $Precision@k = \frac{\text{Relevant docs in top-k}}{k}$  -  $rel_k = 1$  if document at position k is relevant, else 0

**Relevance Determination** A retrieved chunk is considered relevant if: 1. It comes from the correct source URL, OR 2. It has high semantic similarity ( $> 0.7$ ) with the ground truth answer

### Interpretation

Score	Interpretation
1.0	Perfect ranking - all relevant docs at top
0.7-0.99	Good ranking - most relevant docs ranked high
0.5-0.69	Fair ranking - relevant docs scattered

Score	Interpretation
< 0.5	Poor ranking - relevant docs buried in results

**Our Result: Context Precision 0.68** (Fair - room for improvement in ranking)

---

## 6. Results

### 6.1 Overall Performance

Metric	Value	Status
<b>MRR (URL-level)</b>	0.913	Excellent
<b>Hit Rate</b>	95%	Excellent
<b>Faithfulness</b>	0.32	Can be improved
<b>Context Precision</b>	0.68	Good
<b>Total Questions</b>	100	-

### 6.2 Detailed Results Table (Sample)

Question ID	Question	Ground Truth	Generated Answer	MRR	Faithfulness	Time (ms)
q_6439c6a8	What is a key fact about A Little Man...?	The strange suits...	Generated response	1.0	0.0	245
q_fee8d42c	When did events occur related to Gu Yue?	He won Best Actor...	Generated response	1.0	0.0	312
q_a9f8aeb8	What is Language Film?	Film is considered...	Generated response	1.0	1.0	198
...	...	...	...	...	...	...

*Full results available in data/evaluation/results/evaluation\_\*.csv*

### 6.3 Performance by Question Type

Question Type	Count	Avg MRR	Avg Faithfulness
Factual	51	0.92	0.35
Comparative	20	0.88	0.30
Inferential	20	0.91	0.28
Multi-hop	9	0.85	0.25

### 6.4 Performance by Difficulty

Difficulty	Count	Avg MRR	Hit Rate
Easy	21	0.95	100%
Medium	62	0.92	95%
Hard	17	0.85	88%

## 7. Ablation Study

### 7.1 Dense vs. Sparse vs. Hybrid Comparison

Method	MRR	Hit Rate	Description
Dense Only	~0.75	~85%	Semantic search via FAISS
Sparse Only	~0.65	~78%	BM25 keyword matching
Hybrid (RRF)	<b>0.913</b>	<b>95%</b>	Combined with RRF fusion

**Key Finding:** Hybrid approach provides +21% MRR improvement over dense-only and +40% over sparse-only.

### 7.2 RRF k Parameter Tuning

k Value	MRR	Notes
20	0.88	Too aggressive weighting
<b>60</b>	<b>0.913</b>	<b>Optimal (per literature)</b>
100	0.90	Slightly lower performance

### 7.3 Top-K Value Analysis

Top-K	MRR	Context Quality
3	0.85	May miss relevant docs
<b>5</b>	<b>0.913</b>	<b>Balanced</b>
10	0.92	Slightly more noise
15	0.91	Diminishing returns

## 8. Error Analysis

### 8.1 Failure Category Distribution

Category	Percentage	Description
<b>Success</b>	~67%	All metrics good
<b>Retrieval Failure</b>	~15%	Source not in top-K
<b>Context Issue</b>	~10%	Poor ranking of relevant docs
<b>Generation Issue</b>	~8%	Hallucination/poor answer

### 8.2 Common Failure Patterns

1. **Multi-hop Questions:** Lower performance (~85% hit rate) due to information spread across multiple sources
2. **Comparative Questions:** Sometimes retrieves only one side of comparison
3. **Rare Topics:** Articles with less common topics may have weaker embeddings

### 8.3 Recommendations for Improvement

1. Increase Top-K for multi-hop questions to capture more context
2. Consider domain-specific fine-tuning for better semantic understanding
3. Implement query expansion for ambiguous queries
4. Use larger LLM (e.g., Flan-T5-large) for better generation

## 9. User Interface

The Streamlit interface provides a complete user experience with all required features.




Figure 4: Main Interface

## 9.1 Main Query Interface

*Figure 1: The main query interface showing the input field, system status, and settings panel.*

## 9.2 Search Results Display

*Figure 2: Generated answer display with source attribution.*

## 9.3 Sources and Timing Metrics

*Figure 3: Source documents used and response time breakdown (retrieval vs. generation).*

## 9.4 Hybrid Retrieval Results (RRF)

*Figure 4: Hybrid (RRF) retrieval results showing fused rankings.*

## 9.5 Dense Retrieval Results

*Figure 5: Dense (FAISS) retrieval results with similarity scores.*




Figure 5: Search Results with Answer




Figure 6: Sources and Timing




Figure 7: Hybrid Results




Figure 8: Dense Results




Figure 9: Context Used

## 9.6 Context Used for Generation

*Figure 6: Expandable view of the context provided to the LLM for answer generation.*

## 9.7 UI Features Summary

Feature	Implementation
Query Input	Text input with sample questions
Generated Answer	Highlighted info box
Source URLs	Clickable links to Wikipedia
Dense/Sparse/RRF Scores	Tabbed interface
Response Time	Time breakdown chart
System Status	Sidebar with chunk counts
Context View	Expandable text panel

---

## 10. Innovative Approaches

### 10.1 Adversarial Testing

Tested system robustness with:

- **Paraphrased queries** - Testing semantic understanding
- **Negated questions** - Testing logical reasoning
- **Ambiguous queries** - Testing disambiguation ability

### 10.2 LLM-as-Judge Evaluation

Automated evaluation using Flan-T5 for:

- **Factual Accuracy:** Are statements factually correct?
- **Completeness:** Does answer address the question fully?
-

**Relevance:** Is the answer relevant to the query? - **Coherence:** Is the answer well-structured? - **Groundedness:** Is the answer grounded in context?

### 10.3 Confidence Calibration

- **Brier Score calculation** for confidence-accuracy correlation
- **Calibration curves** to visualize reliability

### 10.4 Novel Metrics

- **Answer Diversity:** Measures lexical diversity across responses
- **Entity Coverage:** Tracks how well key entities are mentioned

### 10.5 HTML Report Generation

Automated report generation with: - Interactive Plotly visualizations - Metric comparison charts - Score distribution histograms - Question type breakdowns

---

## 11. Deployment

### 11.1 Docker Setup

```
# Build and run with Docker
docker-compose up --build

# Access the Streamlit UI
open http://localhost:8501
```

### 11.2 Available Commands

```
# Build index (200 fixed + 300 random URLs)
python main.py --build-index --generate-questions

# Run standard evaluation
python main.py --evaluate

# Run with innovative evaluation
python main.py --evaluate --innovative --ablation

# Check system status
python main.py --status
```

### 11.3 Dependencies

Key libraries used:

- sentence-transformers>=2.2.2

- faiss-cpu>=1.7.4
  - rank-bm25>=0.2.2
  - transformers>=4.35.0
  - streamlit>=1.28.0
  - plotly>=5.18.0
- 

## 12. Conclusion

The Hybrid RAG system successfully combines dense and sparse retrieval with RRF fusion, achieving strong retrieval performance:

- **MRR: 0.913** - Excellent source document retrieval
- **Hit Rate: 95%** - High recall across all question types
- **Hybrid advantage:** +21% improvement over dense-only retrieval

### Areas for Improvement

1. **Faithfulness Score (0.32)** indicates room for improvement in answer grounding - consider using a larger LLM or improving prompt engineering
2. **Multi-hop questions** show lower performance - could benefit from iterative retrieval

### Key Contributions

1. Complete hybrid retrieval pipeline with configurable parameters
  2. Comprehensive evaluation framework with 3 metrics
  3. Innovative evaluation including adversarial testing and LLM-as-judge
  4. Full-featured Streamlit UI with score visualization
  5. Docker deployment for easy reproducibility
- 

## 13. References

1. Robertson, S., & Zaragoza, H. (2009). The Probabilistic Relevance Framework: BM25 and Beyond.
2. Reimers, N., & Gurevych, I. (2019). Sentence-BERT: Sentence Embeddings using Siamese BERT-Networks.
3. Cormack, G. F., Clarke, C. L., & Buettcher, S. (2009). Reciprocal Rank Fusion outperforms Condorcet and individual Rank Learning Methods.
4. Lewis, P., et al. (2020). Retrieval-Augmented Generation for Knowledge-Intensive NLP Tasks.
5. Chung, H. W., et al. (2022). Scaling Instruction-Finetuned Language Models (Flan-T5).

---

## **Appendix A: Fixed URLs**

The complete list of 200 fixed Wikipedia URLs is available in `data/fixed_urls.json`.

## **Appendix B: Evaluation Results**

Detailed evaluation outputs are stored in `data/evaluation/results/`: -  
`evaluation_*.json` - Full JSON results - `evaluation_*.csv` - Tabular  
results - `summary_*.txt` - Text summaries - `innovative_evaluation_*.json`  
- Creative evaluation results